

University of Moratuwa
Department of Electronic and Telecommunication Engineering



REPORT

Image Processing & Machine Vision


Supervisors – Dr. Ranga Rodrigo



SAIRISAN.R - 200552V

Question 1



```
t1 = np.linspace(0,50,51).astype('uint8') # 0,50,101
t2 = np.linspace(51,100,0).astype('uint8') #51,200,50
t3 = np.linspace(101,255,100).astype('uint8') #201,255,105
t4 = np.linspace(255,150,0).astype('uint8') # 0,50,101
t5 = np.linspace(151,255,105).astype('uint8') #51,200,50
transform = np.concatenate((t1, t2), axis=0).astype('uint8')
transform = np.concatenate((transform, t3), axis=0).astype('uint8')
transform = np.concatenate((transform, t4), axis=0).astype('uint8')
transform = np.concatenate((transform, t5), axis=0).astype('uint8')
```



Question 2

```
t1 = np.linspace(0,50,51).astype('uint8') # 0,50,101
t2 = np.linspace(51,100,0).astype('uint8') #51,200,50
t3 = np.linspace(101,255,100).astype('uint8') #201,255,105
t4 = np.linspace(255,150,0).astype('uint8') # 0,50,101
t5 = np.linspace(151,255,105).astype('uint8') #51,200,50
transform = np.concatenate((t1, t2), axis=0).astype('uint8')
transform = np.concatenate((transform, t3), axis=0).astype('uint8')
transform = np.concatenate((transform, t4), axis=0).astype('uint8')
transform = np.concatenate((transform, t5), axis=0).astype('uint8')

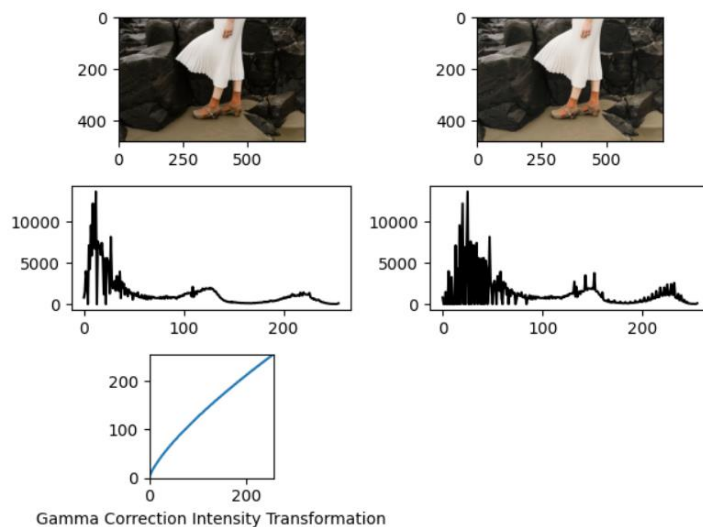
transform=np.flip(transform) #flipping the intensity map array of ealier to get negative.
```



Flip the previous Question transform function to get negative image. So White matter and grey matter are now in their original colour easy to check by medical professionals.

Question 3

```
img_lab = cv.cvtColor(img_orig, cv.COLOR_BGR2LAB) # Convert the image to LAB color space
L_channel = img_lab[:, :, 0] # Extract the L channel
# gamma correction to the L channel
gamma = 0.75 #possible value make image looking good.
table = np.array([(i/255.0)**gamma*255.0 for i in np.arange(0, 256)]).astype('uint8')
L_channel_gamma = cv.LUT(L_channel, table)
img_lab_gamma = img_lab.copy() # Replace the gamma corrected L channel in the LAB image
img_lab_gamma[:, :, 0] = L_channel_gamma
# Convert the LAB image with gamma corrected L channel back to RGB
img_gamma = cv.cvtColor(img_lab_gamma, cv.COLOR_LAB2BGR)
img_gamma = cv.cvtColor(img_gamma, cv.COLOR_BGR2RGB)
```



Gamma value of 0.5 is reasonable value that make good visual to the image. Here gamma should be less than 1 because in the histogram we can visualize that dark pixels (intensity) are higher. And this makes dark areas to light areas slightly.


Question 4

```
hsv_image = cv.cvtColor(image, cv.COLOR_BGR2HSV) #BGR to HSV
hue, saturation, value = cv.split(hsv_image) # Split the HSV image into separate channels
a = 0.5
sigma=40
func=saturation+a*128*np.exp(-((saturation-128)**2)/(2*((sigma)**2)))
transform=np.minimum(func,255).astype('uint8') # intensity transformation
enhanced_hsv_image = cv.merge((hue, transform, value)) # Combine the transformed saturation and value channels
enhanced_image = cv.cvtColor(enhanced_hsv_image, cv.COLOR_HSV2BGR)
```


This code enhances image saturation using a Gaussian-based transformation. The original

img Original Image file
vibra 0 Vibrance-Enhanced Image

```
# defining a function for histogram equalization
def histo_normalized(image):
    hist, bins = np.histogram(image.ravel(), 256, [0, 256])
    cdf = hist.cumsum()
    cdf_normalized = cdf * hist.max() / cdf.max()
    equ = cv.equalizeHist(image)
    return cdf_normalized, equ
#apply hist equalize
cdf_normalized, equalize=histo_normalized(img)
```




Question 5



This code applies histogram equalization to enhance the contrast of a grayscale image. It computes the histogram and cumulative distribution function (CDF) of the original image, then equalizes the image using the CDF. The resulting equalized image has a more balanced distribution of pixel intensities, improving overall contrast.

Question 6



We can clearly see that saturation plane is best to choose for separating foreground helps to create masking image. After that masking image is used to extract foreground of original image and each colour channel histogram are equalized and combined with background of original image. Here background can be generated by inverse of masking image is used as mask to extract background image from original.

```
def resize_img(img,scale):
    new_width = int(img.shape[1] * scale)
    new_height = int(img.shape[0] * scale)
    img=cv.resize(img, (new_width, new_height))
    return img
```

```
threshold,t=cv.threshold(saturation,11,255,cv.THRESH_BINARY)
masked=cv.bitwise_and(image,image,mask=t)
cv.imshow('masked', masked)
cv.imshow('masking image', t) #t=mask image

color = ('b', 'g', 'r')

b, g, r = cv.split(masked)
array=[b,g,r]
array_new=[]

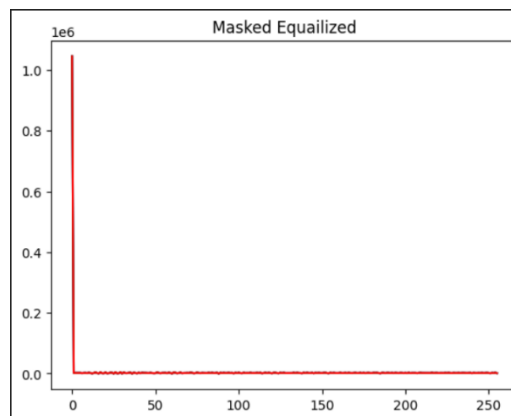

for i, c in enumerate(color):
    hist, bins = np.histogram(masked.ravel(), 256, [0, 256])
    cdf = hist.cumsum()
    cdf_normalized = cdf * hist.max() / cdf.max()
    plt.plot(hist, color=c) #plot the graph for need color
    equi = cv.equalizeHist(array[i])
    array_new.append(equi)

plt.title("Masked")
plt.show()

masked_new = cv.merge(array_new)

for i, c in enumerate(color):
    hist, bins = np.histogram(masked_new.ravel(), 256, [0, 256])
    cdf = hist.cumsum()
    cdf_normalized = cdf * hist.max() / cdf.max()
    plt.plot(hist, color=c) #plot the graph for need color

plt.title("Masked Equailized")
```



Since there are only few changes in output final image because in the masked image the greatest number of pixels are in 0 intensity so histogram equalization would have a minimal impact. Most of the pixel would still remain dark, and the transformation would primarily redistribute intensities within the dark region, making subtle adjustments but not significantly altering the bright pixel.

Question 7

```
kernel1 = np.array([[ -1,  0,  1],
                    [ -2,  0,  2],
                    [ -1,  0,  1]], dtype=np.float32)

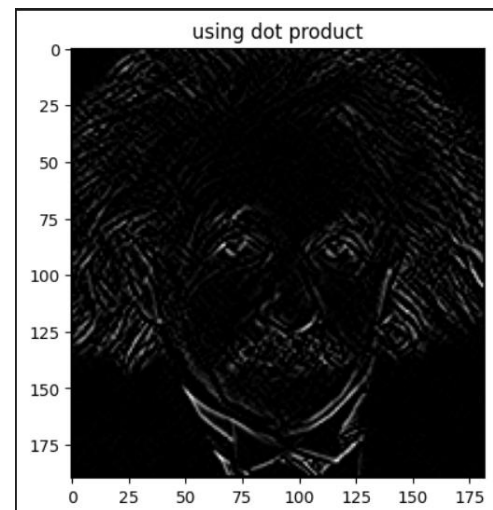

img_filtered1 = cv.filter2D(img, -1, kernel1)

# Display the filtered image
axes[2].imshow(img_filtered1, cmap='gray')
axes[2].set_title('Sobel Horizontal')
axes[2].set_xticks([]), axes[2].set_yticks([])
plt.show()
```

```
# Apply Sobel filter to the image for edge detection
sobel_x = cv.Sobel(img, cv.CV_64F, 1, 0, ksize=5)
sobel_y = cv.Sobel(img, cv.CV_64F, 0, 1, ksize=5)
```

```
# Define the 1D row and column matrices for Sobel filter
sobel_x = np.array([[ -1,  0,  1]]) # Sobel filter for x
sobel_y = np.array([[ -1], [ 0], [ 1]]) # Sobel filter for y

# Create the 2D Sobel filter by multiplying the 1D matrices
sobel_filterx = np.dot(sobel_y, sobel_x)
```



Question 8

```
def zoom_nearest_neighbor(image, scale):
    height, width, channels = image.shape
    new_height = int(height * scale)
    new_width = int(width * scale)

    zoomed_image = np.zeros((new_height, new_width, channels), dtype=np.uint8)

    for i in range(new_height):
        for j in range(new_width):
            original_i = int(i / scale)
            original_j = int(j / scale)
            zoomed_image[i, j] = image[original_i, original_j]

    return zoomed_image
```

SSD Value: 847702528.0

```
def zoom_bilinear(image, scale):
    height, width, channels = image.shape
    new_height = int(height * scale)
    new_width = int(width * scale)

    zoomed_image = np.zeros((new_height, new_width, channels), dtype=np.uint8)

    for i in range(new_height):
        for j in range(new_width):
            # Calculate the corresponding coordinates in the original image
            original_i = i / scale
            original_j = j / scale

            i1, j1 = int(original_i), int(original_j)
            i2, j2 = min(i1 + 1, height - 1), min(j1 + 1, width - 1) #dimension shld not exceed image

            # Calculate the interpolation weights
            difx, dify = original_i - i1, original_j - j1 # 0.67, 0.34
            w1 = (1 - difx) * (1 - dify) # 0.67 * 0.34
            w2 = difx * (1 - dify)
            w3 = (1 - difx) * dify
            w4 = difx * dify

            # Perform bilinear interpolation
            pixel_value = (image[i1, j1] * w1 + image[i1, j2] * w2 +
                           image[i2, j1] * w3 + image[i2, j2] * w4).astype(np.uint8)

            zoomed_image[i, j] = pixel_value

    return zoomed_image
```

SSD Value: 1239645952.0


```
def calculate_ssd(image1, image2):
    diff = image1.astype(np.float32) - image2.astype(np.float32)
    squared_diff = np.square(diff)
    ssd = np.sum(squared_diff)

    return ssd

# zoomed_large = zoom_nearest_neighbor(large_im, 4) #uncomment to see nearest
# zoomed_small=zoom_nearest_neighbor(small_im, 4)
zoomed_large = zoom_bilinear(large_im, 4)
zoomed_small=zoom_bilinear(small_im, 4)
zoomed_large = cv.resize(zoomed_large, (zoomed_small.shape[1], zoomed_small.shape[0]))
```


Question 9

```
mask = np.zeros(image.shape[:2], np.uint8)
bg = np.zeros((1, 65), np.float64)
fg = np.zeros((1, 65), np.float64)
rect = (10, 10, image.shape[1] - 10, image.shape[0] - 10) # Define a rough rectangle around the flower
cv2.grabCut(image, mask, rect, bg, fg, 5, cv2.GC_INIT_WITH_RECT)
mask2 = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')
segmented_image = image * mask2[:, :, np.newaxis]
blurred_background = cv2.GaussianBlur(image, (0, 0), 10)
enhanced_image = image.copy() # Apply the mask for focus
enhanced_image[np.where(mask2[:, :, np.newaxis] == 0)] = blurred_background[np.where(mask2[:, :, np.newaxis] == 0)]
```



The reason the background just beyond the edge of the flower is quite dark in the enhanced image is because, during the blurring step, the pixels in the background are blurred significantly, causing them to become darker and less distinct, while the flower remains sharp and in focus.

Github - [SAIRISAN123/Intensity-Transformations-and-Neighborhood-Filtering \(github.com\)](https://github.com/SAIRISAN123/Intensity-Transformations-and-Neighborhood-Filtering)